

1. Hash tables and hash functions

1.1. Hash tables

1.1.1. Constructors

```
(make-eq-hash-table)
(make-eq-hash-table k)
```

Returns a newly allocated hash table that accepts arbitrary objects as keys, and compares those keys with `eq?`. If an argument is given, then the initial capacity of the hash table is set to *k* elements.

```
(make-eqv-hash-table)
(make-eqv-hash-table k)
```

Returns a newly allocated hash table that accepts arbitrary objects as keys, and compares those keys with `eqv?`. If an argument is given, then the initial capacity of the hash table is set to *k* elements.

```
(make-hash-table procedure1 procedure2)
(make-hash-table procedure1 procedure2 k)
```

Returns a newly allocated mutable hash table using *procedure*₁ as the hash function and *procedure*₂ as the procedure used to compare keys. The hash function must accept a key and return a non-negative exact integer. If a third argument is given, then the initial capacity of the hash table is set to *k* elements.

Both the hash function *procedure*₁ and the comparison predicate *procedure*₂ must behave like pure functions on the domain of keys. For example, the `string-hash` and `string=?` procedures are permissible only if all keys are strings and the contents of those strings are never changed so long as any of them continue to serve as a key in the hash table. Furthermore any pair of values for which the comparison predicate *procedure*₂ returns true must be hashed to the same exact integers by the hash function *procedure*₁.

Note: Hash tables are allowed to cache the results of calling the hash function and comparison predicate, so programs cannot rely on the hash function being called for every lookup or update. Furthermore any hash table operation may call the hash function more than once.

Rationale: Hash table lookups are often followed by updates, so caching may improve performance. Hash tables are free to change their internal representation at any time, which may result in many calls to the hash function.

1.1.2. Procedures

```
(hash-table? hash-table)
```

Returns `#t` if *hash-table* was created by one of the hash table constructors, otherwise returns `#f`.

`(hash-table-size hash-table)`

Returns the number of keys contained in *hash-table* as an exact integer.

`(hash-table-ref hash-table key default)`

Returns the value in *hash-table* associated with *key*. If *hash-table* does not contain an association for *key*, then *default* is returned.

`(hash-table-set! hash-table key value)`

Changes *hash-table* to associate *key* with *value*, replacing any existing association for *key*. Returns the unspecified value.

`(hash-table-delete! hash-table key)`

Removes any association for *key* within *hash-table*. Returns the unspecified value.

`(hash-table-contains? hash-table key)`

Returns `#t` if *hash-table* contains an association for *key*, otherwise returns `#f`.

`(hash-table-update! hash-table key procedure default)`

Equivalent to, but potentially more efficient than:

```
(hash-table-set!  
  hash-table key  
  (procedure (hash-table-ref  
              hash-table key default)))
```

If *hash-table* does not contain an association for *key*, then *default* is passed to *procedure*.

`(hash-table-fold hash-table procedure init)`

For every association in *hash-table*, calls *procedure* with three arguments: the association key, the association value, and an accumulated value. The accumulated value is *init* for the first invocation of *procedure*, and for subsequent invocations of *procedure*, it is the return value of the previous invocation of *procedure*. The order of the calls to *procedure* is indeterminate. The return value of `hash-table-fold` is the value of the last invocation of *procedure*. If any side effect is performed on the hash table while a `hash-table-fold` operation is in progress, then the behavior of `hash-table-fold` is unspecified.

`(hash-table-copy hash-table)`

`(hash-table-copy hash-table immutable)`

Returns a copy of *hash-table*. If the *immutable* argument is provided and is a true value, the returned hash table will be immutable, otherwise it will be mutable.

`(hash-table-clear! hash-table)`

Removes all associations from *hash-table*. Returns the unspecified value.

`(hash-table-for-each procedure hash-table)`

For every association in *hash-table*, calls *procedure* with two arguments: the association key and the association value. The *procedure* is called once for each association in hash-table. The order of these calls is indeterminate. If any side effect is performed on the hash table while a `hash-table-for-each` operation is in progress, then the behavior of `hash-table-for-each` is unspecified. The return value of `hash-table-for-each` is the unspecified value.

Equivalent to:

```
(hash-table-fold hash-table
  (lambda (k v a)
    (procedure k v)
    (unspecified))
  (unspecified))
```

`(hash-table->alist hash-table)`

Equivalent to:

```
(hash-table-fold hash-table
  (lambda (k v a)
    (cons (cons k v) a))
  '())
```

`(hash-table-keys hash-table)`

Equivalent to:

```
(hash-table-fold hash-table
  (lambda (k v a) (cons k a))
  '())
```

`(hash-table-values hash-table)`

Equivalent to:

```
(hash-table-fold hash-table
  (lambda (k v a) (cons v a))
  '())
```

1.1.3. Reflection

(hash-table-equivalence-predicate *hash-table*)

Returns the equivalence predicate used by *hash-table* to compare keys. For hash tables created with `make-eq-hash-table` and `make-equiv-hash-table`, returns `eq?` and `equiv?` respectively.

(hash-table-hash-function *hash-table*)

Returns the hash function used by *hash-table*. For hash tables created by `make-eq-hash-table` or `make-equiv-hash-table`, `#f` is returned.

Rationale: The `make-eq-hash-table` and `make-equiv-hash-table` constructors are designed to hide their hash function. This allows implementations to use the machine address of an object as its hash value, rehashing parts of the table as necessary whenever the garbage collector moves objects to a different address.

(hash-table-mutable? *hash-table*)

Returns `#t` if *hash-table* is mutable, `#f` otherwise.

1.2. Hash functions

The `equal-hash`, `string-hash`, and `string-ci-hash` procedures of this section are acceptable as hash functions only if the keys on which they are called do not suffer side effects while the hash table remains in use.

(equal-hash *obj*)

Returns an integer hash value for *obj*, based on its structure and current contents.

(string-hash *string*)

Returns an integer hash value for *string*, based on its current contents.

(string-ci-hash *string*)

Returns an integer hash value for *string* based on its current contents, ignoring case.

(symbol-hash *symbol*)

Returns an integer hash value for *symbol*.

1.3. Issues

1.3.1. Limit proposal scope

Should the specification be limited to `eq` and `equiv` hash tables, since those are the only kind which may not be implementable as a portable library? (Editors straw-poll: no)

1.3.2. Complexity

It may be appropriate to specify constraints on complexity, such as constant time for `hash-table-size`, or an appropriate constraint on accessor procedures. (Will: more trouble than it's worth in my opinion, because any bound on amortized complexity would have to take into account the cost of rehashing after garbage collections.)

1.3.3. Concurrency

R6RS does not deal with concurrency. Even if this proposal does not say anything about that, the issue should be considered. Any implementation that supports concurrency will have to implement some kind of mutual exclusion for operations that have side effects, and some will need mutual exclusion even for `hash-table-ref`. As specified, the updating operations are not atomic, so they create no new problems. The `hash-table-fold` and `hash-table-for-each` procedures already have a problem, even without concurrency.

1.3.4. `hash-table-update!`

Are the functional update procedures `hash-table-update!` and `hash-table-update!/call` justified? (Will: These may be procedural updates, but they aren't functional. I have flushed `hash-table-update!/call` for the time being, and would happily flush `hash-table-update!` as well.)

1.3.5. Side-effects

There is a potential problem with the higher-order procedures, if a procedure argument mutates the hash table being operated on. This should be addressed somehow, if only by a statement that the behavior caused by such procedures is unspecified.

1.3.6. `hash-table-map`

This has been omitted because a single appropriate specification is not obvious, and any reasonable specification can easily be implemented in terms of `hash-table-fold`.

1.3.7. Names

Should the type name be `hash-table`, `hashtable`, or something else?