

Title

R6RS Library Syntax

Authors

Matthew Flatt and Kent Dybvig

Status

This SRFI is being submitted by a member of the Scheme Language Editor's Committee as part of the R6RS Scheme standardization process. The purpose of such "R6RS SRFIs" is to inform the Scheme community of features and design ideas under consideration by the editors and to allow the community to give the editors some direct feedback that will be considered during the design process.

At the end of the discussion period, this SRFI will be withdrawn. When the R6RS specification is finalized, the SRFI may be revised to conform to the R6RS specification and then resubmitted with the intent to finalize it. This procedure aims to avoid the situation where this SRFI is inconsistent with R6RS. An inconsistency between R6RS and this SRFI could confuse some users. Moreover it could pose implementation problems for R6RS compliant Scheme systems that aim to support this SRFI. Note that departures from the SRFI specification by the Scheme Language Editor's Committee may occur due to other design constraints, such as design consistency with other features that are not under discussion as SRFIs.

1. Abstract

The library system presented here is designed to let programmers share libraries, i.e., code that is intended to be incorporated into larger programs, and especially into programs that use library code from multiple sources. The library system supports macro definitions within libraries, allows macro exports, and distinguishes the phases in which definitions and imports are needed. This SRFI defines a standard notation for libraries, a semantics for library expansion and execution, and a simple format for sharing libraries.

2. Rationale

This standard addresses the following specific goals:

- Separate compilation and analysis; no two libraries have to be compiled at the same time (i.e., the meanings of two libraries cannot depend on each other cyclically, and compilation of two different libraries cannot rely on state shared across compilations), and significant program analysis does not require a whole program.
- Independent compilation/analysis of unrelated libraries, where "unrelated" means that neither depends on the other through a transitive closure of imports.
- Explicit declaration of dependencies, so that the meaning of each identifier is clear at compile time, and so that there is no ambiguity about whether a library needs to be executed for another library's compile time and/or run time.
- Namespace management, so that different library producers are unlikely to define the same top-level name.

It does not address the following:

- Mutually dependent libraries.
- Separation of library interface from library implementation.
- Code outside of a library (e.g., 5 by itself as a program).
- Local modules and local imports.

3. Specification

3.1. Library form

A library declaration contains the following elements:

- a name for the library (possibly compound, with versioning),
- a list of import dependencies, where each dependency specifies the following:
 - the imported library’s name,
 - the relevant phases, e.g., expand or run time, and
 - the subset of the library’s exports to make available within the importing library, and the local names to use within the importing library for each of the library’s exports,
- a list of exports, which name a subset of the library’s imports and definitions, and
- a library body, consisting of a sequence of definitions followed by a sequence of expressions.

3.2. Syntax and Semantics

A library definition is written with the library form:

```
(library <library name>
  (import <import spec>*)
  (export <export spec>*)
  <library body>)
```

The `<library name>` specifies the name of the library, the `import` form specifies the imported bindings, and the `export` form specifies the exported bindings. The `<library body>` specifies the set of definitions, both for local (unexported) and exported bindings, and the set of initialization expressions (commands) to be evaluated for their effects. The exported bindings may be defined within the library or imported into the library. Bindings defined with a library are not visible in code that appears outside of the library unless they are explicitly exported from the library.

No identifier can be imported multiple times, defined multiple times, or both defined and imported.

Library names are consist of a sequence of identifiers in parentheses, followed optionally by a version consisting of a sequence of unsigned integers:

$$\langle \text{library name} \rangle \longrightarrow (\langle \text{identifier} \rangle^+) \mid ((\langle \text{identifier} \rangle^+ \langle \text{version} \rangle))$$
$$\langle \text{version} \rangle \longrightarrow (\langle \text{uinteger } 10 \rangle^+)$$

Each `import spec` specifies a set of bindings to be imported into the library, the phases in which they are to be available, and the local names by which they are to be known.

$$\begin{array}{lcl} \langle \text{import spec} \rangle & \longrightarrow & \langle \text{import set} \rangle \\ & | & (\text{for } \langle \text{import set} \rangle \langle \text{import phase} \rangle^*) \end{array}$$

Valid import phases are **run**, **expand**, **eval**, and **(meta *n*)**, where **run** is an abbreviation for **(meta 0)** and **expand** is an abbreviation for **and (meta 1)**.

$$\langle \text{import phase} \rangle \longrightarrow \text{run} \mid \text{expand} \mid \text{eval} \mid (\text{meta } \langle \text{level} \rangle)$$

$$\langle \text{level} \rangle \longrightarrow \langle \text{uinteger } 10 \rangle$$

Phases are discussed in Section 3.3.

An $\langle \text{import set} \rangle$ names a set of bindings from another library, and possibly specifies local names for the imported bindings.

$$\begin{array}{lcl} \langle \text{import set} \rangle & \longrightarrow & \langle \text{library reference} \rangle \\ & | & (\text{only } \langle \text{import set} \rangle \langle \text{identifier} \rangle^*) \\ & | & (\text{except } \langle \text{import set} \rangle \langle \text{identifier} \rangle^*) \\ & | & (\text{add-prefix } \langle \text{import set} \rangle \langle \text{identifier} \rangle) \\ & | & (\text{rename } \langle \text{import set} \rangle (\langle \text{identifier} \rangle \langle \text{identifier} \rangle)^*) \end{array}$$

A $\langle \text{library reference} \rangle$ identifies a library by its (possibly compound) name and optionally by its version.

$$\langle \text{library reference} \rangle \longrightarrow \langle \text{identifier} \rangle \mid (\langle \text{identifier} \rangle^+) \mid (\langle \text{identifier} \rangle^+ \langle \text{version reference} \rangle)$$

$$\langle \text{version reference} \rangle \longrightarrow (\langle \text{uinteger } 10 \rangle^+) \mid \langle \text{version condition} \rangle$$

$$\begin{array}{lcl} \langle \text{version condition} \rangle & \longrightarrow & (\langle \text{uinteger } 10 \rangle^+) \\ & | & (>= \langle \text{uinteger } 10 \rangle) \\ & | & (<= \langle \text{uinteger } 10 \rangle) \\ & | & (\text{and } \langle \text{version condition} \rangle^+) \\ & | & (\text{or } \langle \text{version condition} \rangle^+) \\ & | & (\text{not } \langle \text{version condition} \rangle) \end{array}$$

The sequence of identifiers in the importing library’s **library reference** must match the sequence of identifiers in the imported library’s **library name**. The importing library’s $\langle \text{version reference} \rangle$ specifies a predicate on a prefix of the imported library’s $\langle \text{version} \rangle$. Each integer must match exactly and each condition has the expected meaning. Everything beyond the prefix specified in the version reference matches unconditionally. When more than one library is identified by a library reference, the choice of libraries is determined in some implementation-dependent manner.

To avoid problems such as incompatible types and replicated state, two libraries whose library names contain the same sequence of identifiers but whose versions do match cannot co-exist in the same program.

By default, all of an imported library’s exported bindings are made visible within an importing library using the names given to the bindings by the imported library. The precise set of bindings to be imported and the names of those bindings can be adjusted with the **only**, **except**, **add-prefix**, and **rename** forms as described below.

- The **only** form produces a subset of the bindings from another $\langle \text{import set} \rangle$, including only the listed $\langle \text{identifier} \rangle$ s; if any of the included $\langle \text{identifier} \rangle$ s is not in $\langle \text{import set} \rangle$, an exception is raised.
- The **except** form produces a subset of the bindings from another $\langle \text{import set} \rangle$, including all but the listed $\langle \text{identifier} \rangle$ s; if any of the excluded $\langle \text{identifier} \rangle$ s is not in $\langle \text{import set} \rangle$, an exception is raised.
- The **add-prefix** adds a prefix to each name from another $\langle \text{import set} \rangle$.
- The **rename** form, for each pair of identifiers $(\langle \text{identifier} \rangle \langle \text{identifier} \rangle)$, removes a binding from the set from $\langle \text{import set} \rangle$, and adds it back with a different name. The first identifier is the original name, and

the second identifier is the new name. If the original name is not in $\langle \text{import set} \rangle$, or if the new name is already in $\langle \text{import set} \rangle$, an exception is raised.

An $\langle \text{export set} \rangle$ names a set of imported and locally defined bindings to be exported, possibly giving them different external names.

$$\begin{aligned} \langle \text{export spec} \rangle &\longrightarrow \langle \text{export set} \rangle \\ &\quad | \quad (\text{for } (\langle \text{export set} \rangle^*) \langle \text{import phase} \rangle^*) \\ \langle \text{export set} \rangle &\longrightarrow \langle \text{identifier} \rangle \\ &\quad | \quad (\text{rename } (\langle \text{identifier} \rangle \langle \text{identifier} \rangle)^*) \end{aligned}$$

In an $\langle \text{export set} \rangle$, an $\langle \text{identifier} \rangle$ names a single binding defined within the library or imported, where the external name for the export is the same as the name of the binding within the library. A **rename** set exports the binding named by the first $\langle \text{identifier} \rangle$ in each pair, using the second $\langle \text{identifier} \rangle$ as the external name.

The $\langle \text{library body} \rangle$ of a **library** form contains definitions for local and exported bindings and initialization expressions to be evaluated when the library is invoked.

A $\langle \text{library body} \rangle$ is like **lambda** body (see below) except that the latter requires the body to contain at least one expression.

$$\begin{aligned} \langle \text{library body} \rangle &\longrightarrow \langle \text{common body} \rangle \\ \langle \text{common body} \rangle &\longrightarrow \langle \text{declaration} \rangle^* \langle \text{definition} \rangle^* \langle \text{command} \rangle^* \\ &\quad | \quad \langle \text{declaration} \rangle^* \langle \text{definition} \rangle^* \langle \text{body begin} \rangle \langle \text{command} \rangle^* \\ \langle \text{body begin} \rangle &\longrightarrow (\text{begin } \langle \text{common body} \rangle) | \langle \text{derived body begin} \rangle \end{aligned}$$

The $\langle \text{body begin} \rangle$ permitted to allow macros to expand into a sequence including both definitions and expressions at the point where the sequence of definitions ends and the sequence of expressions begins. A $\langle \text{derived body begin} \rangle$ is a syntactic abstraction that expands into a $\langle \text{body begin} \rangle$.

Declarations allow a program to specify various *qualities* for the evaluated code. Their interpretation is implementation dependent.

$$\begin{aligned} \langle \text{declaration} \rangle &\longrightarrow (\text{declare } \langle \text{declare spec} \rangle) | \langle \text{derived declaration} \rangle \\ \langle \text{declare spec} \rangle &\longrightarrow \text{unsafe} | \langle \text{quality} \rangle | (\langle \text{quality} \rangle \langle \text{priority} \rangle) \\ \langle \text{quality} \rangle &\longrightarrow \text{safe} | \text{fast} | \text{small} | \text{debug} \\ \langle \text{priority} \rangle &\longrightarrow 0 | 1 | 2 | 3 \end{aligned}$$

A $\langle \text{derived declaration} \rangle$ is a syntactic abstraction that expands into a $\langle \text{declaration} \rangle$.

Definitions consist of syntax definitions, variable definitions, **indirect-export** forms, **begin** forms containing only definitions, and derived definitions.

⟨definition⟩	→	⟨syntax definition⟩
		⟨variable definition⟩
		(indirect-export ⟨indirect spec⟩*)
		(begin ⟨declaration⟩* ⟨definition⟩*)
		⟨derived definition⟩
⟨syntax definition⟩	→	(define-syntax ⟨keyword⟩ ⟨transformer spec⟩)
⟨variable definition⟩	→	(define ⟨variable⟩)
		(define ⟨variable⟩ ⟨expression⟩)
		(define (⟨variable⟩ ⟨def formal⟩) ⟨body⟩)
⟨body⟩	→	⟨common body⟩ ⟨expression⟩

A ⟨derived definition⟩ is a syntactic abstraction that expands into a ⟨definition⟩.

The production for ⟨body⟩ is also relevant for **lambda** bodies, recalling that ⟨lambda expression⟩ is defined as follows.

⟨lambda expression⟩ → (**lambda** ⟨formals⟩ ⟨body⟩)

A ⟨body⟩ is like a ⟨library body⟩ except that at least one expression must be present.

The definitions of a ⟨library body⟩ or ⟨body⟩ are mutually recursive. The transformer expressions and transformer bindings are created from left to right, as described in the R6RS **syntax-case** SRFI [1]. The variable-definition right-hand-side expressions are evaluated from left to right, as if in an implicit **letrec***, and the body expressions are also evaluated from left to right after the variable-definition right-hand-side expressions.

An ⟨indirect spec⟩ declares a set of identifiers to be indirectly exported when a given keyword is exported.

⟨indirect spec⟩ → (⟨keyword⟩ ⟨identifier⟩*)

The specifier declares that the transformer bound to the specified ⟨keyword⟩ may expand into references to the ⟨identifier⟩s and that these ⟨identifier⟩s are to be “indirectly exported” for this purpose if the ⟨keyword⟩ is exported. An exception is raised if any of the specified ⟨keyword⟩s is not defined as a macro within the library, or if any of the associated ⟨identifier⟩s is not defined within the library. An exception is also raised if a macro used outside of its defining library expands into a reference to an identifier that is defined within or imported into the library but is neither directly nor indirectly exported from the library. Identifiers that are not directly exported are not otherwise visible outside of the defining library.

An **indirect-export** may appear anywhere that other definitions may appear, so that syntactic abstractions that expand into syntax definitions can also produce any necessary **indirect-export** forms without concern for the context. They are ignored everywhere except in a ⟨library body⟩.

All exported bindings, explicit or implicit, are immutable both in the exporting and importing libraries.

3.3. Import phases

All bindings imported via a library’s **import** form are *visible* throughout the library’s ⟨library body⟩. An exception is raised, however, if a binding is used out of its declared phase(s):

- Bindings used in run-time code must be imported “for **run**” or “for (**meta** 0).”
- Bindings used in the body of a transformer (appearing on the right-hand-side of a transformer binding) in run-time code must be imported “for **expand**” or “for (**meta** 1),”
- Bindings used in the body of a transformer appearing within the body of a transformer in run-time code must be imported “for (**meta** 2),” and so on.

The import phases of an imported binding are determined by the enclosing `for` form, if any, in the `import` form of the importing library, otherwise the enclosing `for` form, if any, in the `export` form of the exporting library. If neither is given, then the binding is imported “for `run`” only.

The import phases thus implicitly determine when information about a library must be available and also when the various forms contained within a library must be evaluated.

Every library can be characterized by expand-time information (minimally, its imported libraries, a list of the exported keywords, a list of the exported variables, and code to evaluate the transformer expressions) and run-time information (minimally, code to evaluate the variable definition right-hand-side expressions, and code to evaluate the body expressions). The expand-time information must be available to expand references to any exported binding, and the run-time information must be available to evaluate references to any exported variable binding.

If any of a library’s bindings is imported by another library “for `expand`” (or for any meta level greater than 0) both expand-time and run-time information for the first library is made available when the second library is expanded. If any of a library’s bindings is imported by another library “for `run`,” the expand-time information for the first library is made available when the second library is expanded, and the run-time information for the first library is made available when the run-time information for the second library is made available. If any of a library’s bindings is imported by another library “for `eval`,” both expand-time and run-time information for the first library is made available when the run-time information for the second library is made available.

We must also consider when the code to evaluate a library’s transformer expressions is executed and when the code to evaluate the library’s variable-definition right-hand-side expressions and body expressions is executed. We refer to executing the transformer expressions as *visiting* the library and to executing the variable-definition right-hand-side expressions and body expressions as *invoking* the library. A library must be visited before code that uses its bindings can be expanded and invoked before code that uses its bindings can be executed. Visiting or invoking a library may also trigger the visiting or invoking of other libraries.

More precisely, visiting a library at phase N causes the system to:

- Visit at phase N any library that is imported by this library “for `run`” and that is not yet visited at phase N .
- Visit and invoke at phase $N + M$ any library that is imported by this library “for (meta M)” and that is not yet invoked at phase $N + M$.
- Evaluate the library’s transformer expressions.

The order in which imported libraries are visited and invoked is not defined, but imported libraries must be visited and invoked before the library’s transformer expressions are evaluated.

Similarly, invoking a library at meta phase N causes the system to:

- Invoke at phase N any library that is imported by this library “for `run`” and that is not yet invoked at phase N .
- Evaluate the library’s variable-definition right-hand-side and body expressions.

The order in which imported libraries are invoked is not defined, but imported libraries must be invoked before the library’s variable-definition right-hand-side and body expressions are evaluated.

The bindings established by visiting or invoking a library at a given phase are entirely separate from the bindings established by visiting or invoking the library at any other phase, preventing information from being shared through these bindings across phases.

3.4. Eval

The `eval` procedure accepts two arguments, an expression to evaluate, represented as an s-expression, and an environment:

```
(eval expression environment)
```

Environments can be constructed with the `environment` procedure, which accepts a set of import specifiers represented as s-expressions.

```
(environment import-spec ...) ⇒ environment
```

The s-expression syntax of an *import-spec* mirrors the external syntax of an `<import spec>`. For example:

```
(eval '(+ 3 4) (environment '(r6rs))) ⇒ 7
```

An exception is raised if the expand-time or run-time information for a library named in one of the *import-specs* is not *available* when the call to `environment` occurs, in the sense of Section 3.3.

4. Examples

Hello world:

```
(library (hello)
  (import (r6rs))
  (export)
  (display "Hello World")
  (newline))
```

Examples for various `<import spec>`s and `<export spec>`s:

```
(library (stack)
  (import (r6rs))
  (export make push! pop! empty!)

  (define (make) (list '()))
  (define (push! s v) (set-car! s (cons v (car s))))
  (define (pop! s) (let ([v (caar s)])
                    (set-car! s (cдар s))
                    v))
  (define (empty! s) (set-car! s '())))

(library (balloons)
  (import (r6rs))
  (export make push pop)

  (define (make w h) (cons w h))
  (define (push b amt) (cons (- (car b) amt) (+ (cdr b) amt)))
  (define (pop b) (display "Boom! ")
                  (display (* (car b) (cdr b)))
                  (newline)))

(library (party)
  (import (r6rs))
  (only (stack) make push! pop!) ; not empty!
  (add-prefix (balloons) balloon:))
```

```

;; Total exports: make, push, push!, make-party, pop!
(export (rename (balloon:make make)
                (balloon:push push))
push!
make-party
(rename (party-pop! pop!)))

;; Creates a party as a stack of balloons, starting with
;; two balloons
(define (make-party)
  (let ([s (make)]) ; from stack
    (push! s (balloon:make 10 10))
    (push! s (balloon:make 12 9))
    s))
(define (party-pop! p)
  (balloon:pop (pop! p)))

(library (main)
  (import (r6rs) (party))

  (define p (make-party))
  (pop! p)      ; displays "Boom! 108"
  (push! p (push (make 5 5) 1))
  (pop! p)      ; displays "Boom! 24"

```

Examples for macros and phases:

```

(library (helper)
  (import (r6rs))
  (export find-dup)

  (define (find-dup l)
    (and (pair? l)
         (let loop ((rest (cdr l)))
           (cond
            [(null? rest) (find-dup (cdr l))]
            [(bound-identifier=? (car l) (car rest)) (car rest)]
            [else (loop (cdr rest))])))))

(library (mvlet)
  (import (r6rs) (import (for (helper) expand))
  (export mvlet)

  (define-syntax mvlet
    (lambda (stx)
      (syntax-case stx ()
        [(~ [(id ...) expr] body0 body ...)
         (not (find-dup (syntax-object->list (syntax (id ...)))))
         (syntax (call-with-values (lambda () expr)
                                   (lambda (id ...) body0 body ...))))]))))

(library (let-div)
  (import (r6rs) (mvlet))
  (export let-div)

```



```

(define (quotient+remainder n d)
  (let ([q (quotient n d)])
    (values q (- n (* q d)))))
(indirect-export (let-div quotient+remainder))
(define-syntax let-div
  (syntax-rules ()
    [(_ n d (q r) body0 body ...)
     (mvlet [(q r) (quotient+remainder n d)]
       body0 body ...)]))

```

5. Reference Implementation

6. Issues

6.1. Simpler library names

A single identifier, e.g., `r6rs`, could be allowed as an abbreviation for a library name sequence containing just one identifier, e.g., `(r6rs)`.

6.2. Expanding import and export form

The import and export forms could be allowed in macro expansions. In that case, the exports and imports of a library can be determined only by macro-expanding the library.

6.3. co-export form

Better support for records may be useful, possibly through a co-export form: `(co-export <co spec>*)`, where `<co spec>` contains an identifier defined or imported into a library followed by a list of other identifiers defined or imported into the library:

$$\langle \text{co spec} \rangle \longrightarrow (\langle \text{identifier} \rangle \langle \text{identifier} \rangle^*)$$

The meaning of a `<co spec>` is that if the first identifier becomes exported, either through `export` or `co-export`, then the remaining identifiers are also exported. Similarly, if the first identifier is indirectly exported through `indirect-export` or `co-export`, then the remaining identifiers are indirectly exported.

To be useful for managing records, `co-export` must be allowed in macro expansions. In that case, the full exports of a library can be determined only by macro-expanding the library.

6.4. Implicit phasing

Instead of having the programmer declare import phases explicitly, phases could be inferred from uses of the imported identifiers. In that case, determining the phase of an import would require macro-expanding the library body.

7. Acknowledgments

This SRFI was written in consultation with the full set of R6RS editors: Will Clinger, Kent Dybvig, Marc Feeley, Matthew Flatt, Manuel Serrano, Michael Sperber, and Anton van Straaten.

8. References

- [1] R. Kent Dybvig. SRFI 93: R6rs syntax-case macros.
<http://srfi.schemers.org/srfi-93/srfi-93.html>.
- [2] Matthew Flatt. Composable and compilable macros: You want it when? In *International Conference on Functional Programming (ICFP'2002)*, 2002.
<http://www.cs.utah.edu/plt/publications/macromod.pdf>.
- [3] Richard A. Kelsey and Jonathan A. Rees. Modules (in 'the incomplete scheme48 reference manual for release 0.57'). 2001.
http://s48.org/0.57/manual/s48manual_24.html.
- [4] Oscar Waddell and R. Kent Dybvig. Extending the scope of syntactic abstraction. In *Conference Record of the Twenty Sixth Annual ACM Symposium on Principles of Programming Languages*, 203–213, January 1999.

9. Copyright

Copyright © Matthew Flatt and Kent Dybvig (2005–2006). All Rights Reserved.

Permission is hereby granted, free of charge, to any person obtaining a copy of this software and associated documentation files (the "Software"), to deal in the Software without restriction, including without limitation the rights to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the Software, and to permit persons to whom the Software is furnished to do so, subject to the following conditions:

The above copyright notice and this permission notice shall be included in all copies or substantial portions of the Software.

THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE SOFTWARE.